

# TensorFlow: Blas GEMM launch failed

Asked 3 years, 8 months ago   Active 2 days ago   Viewed 35k times

When I'm trying to use TensorFlow with Keras using the gpu, I'm getting this error message:

31

```
C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\ipykernel\__main__.py:2:
UserWarning: Update your `fit_generator` call to the Keras 2 API:
`fit_generator(<keras.pre..., 37800, epochs=2, validation_data=<keras.pre...,
validation_steps=4200)`
from ipykernel import kernelapp as app
```

11

Epoch 1/2

```
InternalError                                Traceback (most recent call last)
C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\client\session.py in _do_call(self, fn, *args)
    1038         try:
-> 1039             return fn(*args)
    1040         except errors.OpError as e:

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\client\session.py in _run_fn(session, feed_dict, fetch_list,
target_list, options, run_metadata)
    1020                                     feed_dict, fetch_list, target_list,
-> 1021                                     status, run_metadata)
    1022
```

```
C:\Users\nicol\Anaconda3\envs\tensorflow\lib\contextlib.py in __exit__(self, type,
value, traceback)
    65         try:
--> 66             next(self.gen)
    67         except StopIteration:
```

```
C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\framework\errors_impl.py in
raise_exception_on_not_ok_status()
    465         compat.as_text(pywrap_tensorflow.TF_Message(status)),
-> 466         pywrap_tensorflow.TF_GetCode(status))
    467     finally:
```

```
InternalError: Blas GEMM launch failed : a.shape=(64, 784), b.shape=(784, 10), m=64,
n=10, k=784
[[Node: dense_1/MatMul = MatMul[T=DT_FLOAT, transpose_a=false, transpose_b=false,
_device="/job:localhost/replica:0/task:0/gpu:0"]](flatten_1/Reshape,
dense_1/kernel/read)]]
```

During handling of the above exception, another exception occurred:

```
InternalError                                Traceback (most recent call last)
<ipython-input-13-2a52d1079a66> in <module>()
      1 history=model.fit_generator(batches, batches.n, nb_epoch=2,
----> 2                             validation_data=val_batches, nb_val_samples=val_batches.n)

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\legacy\interfaces.py
in wrapper(*args, **kwargs)
    86         warnings.warn('Update your `'+ object_name +
    87                        '` call to the Keras 2 API: ' + signature,
stacklevel=2)
--> 88         return func(*args, **kwargs)
```

```

189         wrapper._legacy_support_signature = inspect.getargspec(func)
190         return wrapper

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\models.py in
fit_generator(self, generator, steps_per_epoch, epochs, verbose, callbacks,
validation_data, validation_steps, class_weight, max_q_size, workers, pickle_safe,
initial_epoch)
1108             workers=workers,
1109             pickle_safe=pickle_safe,
-> 1110             initial_epoch=initial_epoch)
1111
1112     @interfaces.legacy_generator_methods_support

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\legacy\interfaces.py
in wrapper(*args, **kwargs)
86         warnings.warn('Update your ``' + object_name +
87                        `` call to the Keras 2 API: ' + signature,
stacklevel=2)
--> 88         return func(*args, **kwargs)
89         wrapper._legacy_support_signature = inspect.getargspec(func)
90         return wrapper

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\engine\training.py in
fit_generator(self, generator, steps_per_epoch, epochs, verbose, callbacks,
validation_data, validation_steps, class_weight, max_q_size, workers, pickle_safe,
initial_epoch)
1888             outs = self.train_on_batch(x, y,
1889                                         sample_weight=sample_weight,
-> 1890                                         class_weight=class_weight)
1891
1892             if not isinstance(outs, list):

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\engine\training.py in
train_on_batch(self, x, y, sample_weight, class_weight)
1631         ins = x + y + sample_weights
1632         self._make_train_function()
-> 1633         outputs = self.train_function(ins)
1634         if len(outputs) == 1:
1635             return outputs[0]

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\backend\tensorflow_backend.py in __call__(self, inputs)
2227         session = get_session()
2228         updated = session.run(self.outputs + [self.updates_op],
-> 2229                             feed_dict=feed_dict)
2230         return updated[:len(self.outputs)]
2231

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\tensorflow\python\client\session.py in run(self, fetches, feed_dict, options,
run_metadata)
776     try:
777         result = self._run(None, fetches, feed_dict, options_ptr,
--> 778                             run_metadata_ptr)
779         if run_metadata:
780             proto_data = tf_session.TF_GetBuffer(run_metadata_ptr)

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\tensorflow\python\client\session.py in _run(self, handle, fetches, feed_dict,
options, run_metadata)
980     if final_fetches or final_targets:
981         results = self._do_run(handle, final_targets, final_fetches,
--> 982                               feed_dict_string, options, run_metadata)
983     else:
984         results = []

```

```

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\client\session.py in _do_run(self, handle, target_list,
fetch_list, feed_dict, options, run_metadata)
    1030     if handle is None:
    1031         return self._do_call(_run_fn, self._session, feed_dict, fetch_list,
-> 1032                             target_list, options, run_metadata)
    1033     else:
    1034         return self._do_call(_prun_fn, self._session, handle, feed_dict,

```

```

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\client\session.py in _do_call(self, fn, *args)
    1050     except KeyError:
    1051         pass
-> 1052     raise type(e)(node_def, op, message)
    1053
    1054     def _extend_graph(self):

```

```

InternalError: Blas GEMM launch failed : a.shape=(64, 784), b.shape=(784, 10), m=64,
n=10, k=784
[[Node: dense_1/MatMul = MatMul[T=DT_FLOAT, transpose_a=false, transpose_b=false,
_device="/job:localhost/replica:0/task:0/gpu:0"]](flatten_1/Reshape,
dense_1/kernel/read)]]

```

Caused by op 'dense\_1/MatMul', defined at:

```

File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\runpy.py", line 193, in
_run_module_as_main
    "__main__", mod_spec)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\runpy.py", line 85, in _run_code
    exec(code, run_globals)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\__main__.py", line 3, in <module>
    app.launch_new_instance()
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\traitlets\config\application.py", line 658, in launch_instance
    app.start()
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\kernelapp.py", line 477, in start
    ioloop.IOLoop.instance().start()
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\zmq\eventloop\ioloop.py", line 177, in start
    super(ZMQIOLoop, self).start()
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\tornado\ioloop.py",
line 888, in start
    handler_func(fd_obj, events)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tornado\stack_context.py", line 277, in null_wrapper
    return fn(*args, **kwargs)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\zmq\eventloop\zmqstream.py", line 440, in _handle_events
    self._handle_recv()
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\zmq\eventloop\zmqstream.py", line 472, in _handle_recv
    self._run_callback(callback, msg)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\zmq\eventloop\zmqstream.py", line 414, in _run_callback
    callback(*args, **kwargs)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tornado\stack_context.py", line 277, in null_wrapper
    return fn(*args, **kwargs)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\kernelbase.py", line 283, in dispatcher
    return self.dispatch_shell(stream, msg)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\kernelbase.py", line 235, in dispatch_shell

```

```

    handler(stream, idents, msg)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\kernelbase.py", line 399, in execute_request
    user_expressions, allow_stdin)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\ipkernel.py", line 196, in do_execute
    res = shell.run_cell(code, store_history=store_history, silent=silent)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\zmqshell.py", line 533, in run_cell
    return super(ZMQInteractiveShell, self).run_cell(*args, **kwargs)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\IPython\core\interactiveshell.py", line 2683, in run_cell
    interactivity=interactivity, compiler=compiler, result=result)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\IPython\core\interactiveshell.py", line 2787, in run_ast_nodes
    if self.run_code(code, result):
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\IPython\core\interactiveshell.py", line 2847, in run_code
    exec(code_obj, self.user_global_ns, self.user_ns)
File "<ipython-input-10-1e7a3b259f23>", line 4, in <module>
    model.add(Dense(10, activation='softmax'))
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\models.py",
line 466, in add
    output_tensor = layer(self.outputs[0])
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\keras\engine\topology.py", line 585, in __call__
    output = self.call(inputs, **kwargs)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\keras\layers\core.py", line 840, in call
    output = K.dot(inputs, self.kernel)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\keras\backend\tensorflow_backend.py", line 936, in dot
    out = tf.matmul(x, y)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\ops\math_ops.py", line 1801, in matmul
    a, b, transpose_a=transpose_a, transpose_b=transpose_b, name=name)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\ops\gen_math_ops.py", line 1263, in _mat_mul
    transpose_b=transpose_b, name=name)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\framework\op_def_library.py", line 768, in apply_op
    op_def=op_def)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\framework\ops.py", line 2336, in create_op
    original_op=self._default_original_op, op_def=op_def)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\framework\ops.py", line 1228, in __init__
    self._traceback = _extract_stack()

InternalError (see above for traceback): Blas GEMM launch failed : a.shape=(64, 784),
b.shape=(784, 10), m=64, n=10, k=784
[[Node: dense_1/MatMul = MatMul[T=DT_FLOAT, transpose_a=false, transpose_b=false,
_device="/job:localhost/replica:0/task:0/gpu:0"]](flatten_1/Reshape,
dense_1/kernel/read)]]

```

When I'm trying to use TensorFlow with Keras using the cpu, I'm getting this error message:

```

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\ipykernel\__main__.py:5:
UserWarning: Update your `fit_generator` call to the Keras 2 API:
`fit_generator(<keras.pre..., 37800, validation_steps=4200, validation_data=
<keras.pre..., epochs=2)`
Epoch 1/2

```

```

-----
InternalError                                Traceback (most recent call last)
C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\client\session.py in _do_call(self, fn, *args)
    1038     try:
-> 1039         return fn(*args)
    1040     except errors.OpError as e:

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\client\session.py in _run_fn(session, feed_dict, fetch_list,
target_list, options, run_metadata)
    1020         feed_dict, fetch_list, target_list,
-> 1021         status, run_metadata)
    1022

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\contextlib.py in __exit__(self, type,
value, traceback)
     65         try:
---> 66             next(self.gen)
     67         except StopIteration:

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\framework\errors_impl.py in
raise_exception_on_not_ok_status()
    465         compat.as_text(pywrap_tensorflow.TF_Message(status)),
-> 466         pywrap_tensorflow.TF_GetCode(status))
    467     finally:

InternalError: Blas GEMM launch failed : a.shape=(64, 784), b.shape=(784, 10), m=64,
n=10, k=784
[[Node: dense_1/MatMul = MatMul[T=DT_FLOAT, transpose_a=false, transpose_b=false,
_device="/job:localhost/replica:0/task:0/gpu:0"]](flatten_1/Reshape,
dense_1/kernel/read)]
[[Node: Assign_3/_84 = _Recv[client_terminated=false,
recv_device="/job:localhost/replica:0/task:0/cpu:0",
send_device="/job:localhost/replica:0/task:0/gpu:0", send_device_incarnation=1,
tensor_name="edge_374_Assign_3", tensor_type=DT_FLOAT,
_device="/job:localhost/replica:0/task:0/cpu:0"]()]]

During handling of the above exception, another exception occurred:

InternalError                                Traceback (most recent call last)
<ipython-input-14-f66b4d3d5b88> in <module>()
      3 with tf.device('/cpu:0'):
      4     history=model.fit_generator(batches, batches.n, nb_epoch=2,
---> 5         validation_data=val_batches, nb_val_samples=val_batches.n)

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\legacy\interfaces.py
in wrapper(*args, **kwargs)
     86         warnings.warn('Update your ` ' + object_name +
     87             '` call to the Keras 2 API: ' + signature,
stacklevel=2)
---> 88         return func(*args, **kwargs)
     89     wrapper._legacy_support_signature = inspect.getargspec(func)
     90     return wrapper

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\models.py in
fit_generator(self, generator, steps_per_epoch, epochs, verbose, callbacks,
validation_data, validation_steps, class_weight, max_q_size, workers, pickle_safe,
initial_epoch)
    1108         workers=workers,
    1109         pickle_safe=pickle_safe,
-> 1110         initial_epoch=initial_epoch)
    1111
    1112     @interfaces.legacy_generator_methods_support

```

```

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\legacy\interfaces.py
in wrapper(*args, **kwargs)
    86         warnings.warn('Update your ` ` + object_name +
    87                        `` call to the Keras 2 API: ' + signature,
stacklevel=2)
--> 88         return func(*args, **kwargs)
    89         wrapper._legacy_support_signature = inspect.getargspec(func)
    90         return wrapper

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\engine\training.py in
fit_generator(self, generator, steps_per_epoch, epochs, verbose, callbacks,
validation_data, validation_steps, class_weight, max_q_size, workers, pickle_safe,
initial_epoch)
    1888             outs = self.train_on_batch(x, y,
    1889                                         sample_weight=sample_weight,
-> 1890                                         class_weight=class_weight)
    1891
    1892             if not isinstance(outs, list):

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\engine\training.py in
train_on_batch(self, x, y, sample_weight, class_weight)
    1631         ins = x + y + sample_weights
    1632         self._make_train_function()
-> 1633         outputs = self.train_function(ins)
    1634         if len(outputs) == 1:
    1635             return outputs[0]

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\keras\backend\tensorflow_backend.py in __call__(self, inputs)
    2227         session = get_session()
    2228         updated = session.run(self.outputs + [self.updates_op],
-> 2229                             feed_dict=feed_dict)
    2230         return updated[:len(self.outputs)]
    2231

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\client\session.py in run(self, fetches, feed_dict, options,
run_metadata)
    776         try:
    777             result = self._run(None, fetches, feed_dict, options_ptr,
--> 778                             run_metadata_ptr)
    779         if run_metadata:
    780             proto_data = tf_session.TF_GetBuffer(run_metadata_ptr)

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\client\session.py in _run(self, handle, fetches, feed_dict,
options, run_metadata)
    980         if final_fetches or final_targets:
    981             results = self._do_run(handle, final_targets, final_fetches,
--> 982                                   feed_dict_string, options, run_metadata)
    983         else:
    984             results = []

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\client\session.py in _do_run(self, handle, target_list,
fetch_list, feed_dict, options, run_metadata)
    1030         if handle is None:
    1031             return self._do_call(_run_fn, self._session, feed_dict, fetch_list,
-> 1032                                   target_list, options, run_metadata)
    1033         else:
    1034             return self._do_call(_prun_fn, self._session, handle, feed_dict,

C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\client\session.py in _do_call(self, fn, *args)

```



```

1050         except KeyError:
1051             pass
-> 1052         raise type(e)(node_def, op, message)
1053
1054     def _extend_graph(self):

```

```

InternalError: Blas GEMM launch failed : a.shape=(64, 784), b.shape=(784, 10), m=64,
n=10, k=784
[[Node: dense_1/MatMul = MatMul[T=DT_FLOAT, transpose_a=false, transpose_b=false,
_device="/job:localhost/replica:0/task:0/gpu:0"] (flatten_1/Reshape,
dense_1/kernel/read)]]
[[Node: Assign_3/_84 = _Recv[client_terminated=false,
recv_device="/job:localhost/replica:0/task:0/cpu:0",
send_device="/job:localhost/replica:0/task:0/gpu:0", send_device_incarnation=1,
tensor_name="edge_374_Assign_3", tensor_type=DT_FLOAT,
_device="/job:localhost/replica:0/task:0/cpu:0"]()]]

```

Caused by op 'dense\_1/MatMul', defined at:

```

File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\runpy.py", line 193, in
_run_module_as_main
    "__main__", mod_spec)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\runpy.py", line 85, in _run_code
    exec(code, run_globals)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\__main__.py", line 3, in <module>
    app.launch_new_instance()
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\traitlets\config\application.py", line 658, in launch_instance
    app.start()
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\kernelapp.py", line 477, in start
    ioloop.IOLoop.instance().start()
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\zmq\eventloop\ioloop.py", line 177, in start
    super(ZMQIOLoop, self).start()
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\tornado\ioloop.py",
line 888, in start
    handler_func(fd_obj, events)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tornado\stack_context.py", line 277, in null_wrapper
    return fn(*args, **kwargs)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\zmq\eventloop\zmqstream.py", line 440, in _handle_events
    self._handle_recv()
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\zmq\eventloop\zmqstream.py", line 472, in _handle_recv
    self._run_callback(callback, msg)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\zmq\eventloop\zmqstream.py", line 414, in _run_callback
    callback(*args, **kwargs)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tornado\stack_context.py", line 277, in null_wrapper
    return fn(*args, **kwargs)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\kernelbase.py", line 283, in dispatcher
    return self.dispatch_shell(stream, msg)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\kernelbase.py", line 235, in dispatch_shell
    handler(stream, idents, msg)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\kernelbase.py", line 399, in execute_request
    user_expressions, allow_stdin)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\ipkernel.py", line 196, in do_execute
    res = shell.run_cell(code, store_history=store_history, silent=silent)

```

```

File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\ipykernel\zmqshell.py", line 533, in run_cell
    return super(ZMQInteractiveShell, self).run_cell(*args, **kwargs)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\IPython\core\interactiveshell.py", line 2683, in run_cell
    interactivity=interactivity, compiler=compiler, result=result)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\IPython\core\interactiveshell.py", line 2787, in run_ast_nodes
    if self.run_code(code, result):
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\IPython\core\interactiveshell.py", line 2847, in run_code
    exec(code_obj, self.user_global_ns, self.user_ns)
File "<ipython-input-12-1e7a3b259f23>", line 4, in <module>
    model.add(Dense(10, activation='softmax'))
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-packages\keras\models.py",
line 466, in add
    output_tensor = layer(self.outputs[0])
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\keras\engine\topology.py", line 585, in __call__
    output = self.call(inputs, **kwargs)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\keras\layers\core.py", line 840, in call
    output = K.dot(inputs, self.kernel)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\keras\backend\tensorflow_backend.py", line 936, in dot
    out = tf.matmul(x, y)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\ops\math_ops.py", line 1801, in matmul
    a, b, transpose_a=transpose_a, transpose_b=transpose_b, name=name)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\ops\gen_math_ops.py", line 1263, in _mat_mul
    transpose_b=transpose_b, name=name)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\framework\op_def_library.py", line 768, in apply_op
    op_def=op_def)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\framework\ops.py", line 2336, in create_op
    original_op=self._default_original_op, op_def=op_def)
File "C:\Users\nicol\Anaconda3\envs\tensorflow\lib\site-
packages\tensorflow\python\framework\ops.py", line 1228, in __init__
    self._traceback = _extract_stack()

InternalError (see above for traceback): Blas GEMM launch failed : a.shape=(64, 784),
b.shape=(784, 10), m=64, n=10, k=784
[[Node: dense_1/MatMul = MatMul[T=DT_FLOAT, transpose_a=false, transpose_b=false,
_device="/job:localhost/replica:0/task:0/gpu:0"]](flatten_1/Reshape,
dense_1/kernel/read)]]
[[Node: Assign_3/_84 = _Recv[client_terminated=false,
recv_device="/job:localhost/replica:0/task:0/cpu:0",
send_device="/job:localhost/replica:0/task:0/gpu:0", send_device_incarnation=1,
tensor_name="edge_374_Assign_3", tensor_type=DT_FLOAT,
_device="/job:localhost/replica:0/task:0/cpu:0"]()]]

```

In both cases, the error is with InternalError (see above for traceback): Blas GEMM launch failed  
Can you tell me how to get Blas GEMM to launch? I installed tensorflow and keras in a 3.5 python anaconda environment where I also installed all needed module (numpy, pandas, scipy, scikit-learn). I have a Windows 10 with a NVIDIA gpu that can use CUDA. I downloaded CUDA and cuDNN. I'm using the Jupyter notebook on Chrome.



Sometimes when I run my code, rather than having this error, I get that it starts running and then it crashes. After the crash, I can't do anything on my jupyter notebook and after some time a pop-up asks me if I want to kill the page. This is an image of what I got after the crash. !

(<http://www.hostingpics.net/viewer.php?id=647186tensorflowError.png>)

P.S. I know my problem is similar as in this question: [Tensorflow Basic Example Error: CUBLAS\\_STATUS\\_NOT\\_INITIALIZED](#) but it has not been solved there and I'm not sure this question is clear enough or is exactly the same problem as I have so I'm posting it with my own error message. This problem is different of: [TensorFlow: InternalError: Blas SGEMM launch failed](#) Since I have a problem with GEMM rather than SGEMM and that my problem is both with gpu and cpu and it is not solved by the answer of this question.

python tensorflow keras blas

Share Edit Follow

edited May 16 '17 at 4:39



talonmies

66.4k 18 163 241

asked May 15 '17 at 22:59



Nicolas

359 1 3 9

## 18 Answers

Active	Oldest	Votes
--------	--------	-------



41



This worked for me on TensorFlow 2.1.0 (per:

[https://www.tensorflow.org/api\\_docs/python/tf/config/experimental/set\\_memory\\_growth](https://www.tensorflow.org/api_docs/python/tf/config/experimental/set_memory_growth))

```
import tensorflow as tf
physical_devices = tf.config.list_physical_devices('GPU')
tf.config.experimental.set_memory_growth(physical_devices[0], True)
```

Share Edit Follow

answered Mar 16 '20 at 1:36



socasanta

871 5 13

3 Thanks, this seems to work for me aswell on Tensorflow 2.1.0 – [charel-f](#) Apr 4 '20 at 10:41

1 What about tensorflow 2.0, if I may ask – [MasterControlProgram](#) May 2 '20 at 17:02

1 @MasterControlProgram It's present in the 2.0 API, but I haven't actually tried it in 2.0. – [socasanta](#) May 3 '20 at 16:00

2 I had to change the code little for a multi-GPU machine but found it eventually worked. – [Tae-Sung Shin](#) May 28 '20 at 18:01



It's a simple fix, but it was a nightmare to figure it all out

24

On Windows I found the Keras install in Anaconda3\Lib\site-packages\keras

sources:

[https://www.tensorflow.org/guide/using\\_gpu](https://www.tensorflow.org/guide/using_gpu)[https://github.com/keras-team/keras/blob/master/keras/backend/tensorflow\\_backend.py](https://github.com/keras-team/keras/blob/master/keras/backend/tensorflow_backend.py)

Find the following in your keras/tensorflow\_backend.py file you'll add  
 config.gpu\_options.allow\_growth= True in both places

```
if _SESSION is None:
    if not os.environ.get('OMP_NUM_THREADS'):
        config = tf.ConfigProto(allow_soft_placement=True)
        config.gpu_options.allow_growth=True
    else:
        num_thread = int(os.environ.get('OMP_NUM_THREADS'))
        config = tf.ConfigProto(intra_op_parallelism_threads=num_thread,
                                allow_soft_placement=True)
        config.gpu_options.allow_growth=True
    _SESSION = tf.Session(config=config)
session = _SESSION
```

Share Edit Follow

answered Sep 1 '18 at 21:39

**Linda MacPhee-Cobb****5,833** 3 16 18

It would be nice to read one line why :) For me, the solution was that the GPU is not used in the background – [MasterControlProgram](#) May 2 '20 at 16:49

2 the above lines are not present in tensorflow 2.0 – [Hrushu](#) May 18 '20 at 9:35



5

Make sure you have no other processes using the GPU running. Run nvidia-smi to check this.

**SOURCE:** [An issue brought up by @reedwm](#).

Share Edit Follow

edited Jun 13 '19 at 17:43

**LogicalBranch****3,392** 2 16 47

answered Jun 13 '19 at 12:48

**liwy****59** 1 2

1 For example when you use pycharm and jupyter in parallel! – [Phil](#) Jul 2 '20 at 14:12

Thanks for this. Previous answers worked for me as well, but I was wondering why the issue was appearing out of no where. After seeing this I killed the running process and the problem was solved automatically. – [Pritam](#) Oct 3 '20 at 18:03

Another example: having multiple kernels running in JupyterLab. Just restarting the kernel with this bug was insufficient; I had to shut down all other kernels first. – [Denziloe](#) Oct 26 '20 at 21:55

This Answer is much related to **Tensorflow**:

5 Sometimes Tensorflow fails at creation in Windows.

***Restarting the notebook using gpu solves it in most cases***

If it doesn't then try restarting the notebook after adding these options in your code.

```
gpu_options = tf.GPUOptions(per_process_gpu_memory_fraction=0.9)
tf.Session(config=tf.ConfigProto(gpu_options=gpu_options, allow_soft_placement=True))
```

I never had such error while using **Keras** But try restarting your notebook

Share Edit Follow

edited Jul 27 '19 at 7:44

answered Sep 13 '18 at 16:10



Michael Yadidya

971 7 15

there is a closing parenthesis missing at the end – [bit\\_scientist](#) Mar 17 '20 at 1:46

Would be a helpful .. Unfortunately it does not have GPUOptions() in the next version 2. Could you update?  
Thanks you! – [MasterControlProgram](#) May 2 '20 at 16:57

Had the same error. Maybe it is related to the problem that tensorflow is [allocating all gpu memory](#). But the fix recommended there didn't work for me and it is not possible yet to limit tensorflow's gpu memory use via keras.json or commandline. Switching keras' backend to Theano resolved the issue for me (howto can be found [here](#)).

Share Edit Follow

edited Jun 16 '17 at 7:57

answered Jun 4 '17 at 12:31



MiniQuark

39k 27 133 164



fotis j

287 4 12

I was getting exactly the same error message. I realized that there was an error with my CUDA installation, specifically with the `cuBLAS` library.

You can check if yours has the same problem by running the sample program `simpleCUBLAS` (it comes with the CUDA installation, you will probably find it in the CUDA home folder: `$CUDA_HOME\samples\7_CUDALibraries\simpleCUBLAS` )

Try running this program. If the test fails, you have a problem with your CUDA installation. You should try to reinstall it. That's how I solved the same problem here.

Renaming `cublas64_10.dll` to `cublas64_100.dll` may be a solution.

Share Edit Follow

edited May 18 '20 at 19:15

answered Jun 16 '17 at 22:07



Nicolas Gervais

17.6k 8 47 76



Tairone

91 1 6

- 5 In case anyone stumbles upon this: Renaming **cublas64\_10.dll** to **cublas64\_100.dll** worked for me. – [asymmetryFan](#) Dec 23 '19 at 14:48



I ran into this problem when trying to run several servers that use a model to serve predictions. As I wasn't training a model but simply using it, the difference between using GPU or CPU was minor. For this specific case, the issue can be avoided by forcing Tensorflow to use the CPU by "hiding" the GPU.



```
import os
os.environ["CUDA_VISIBLE_DEVICES"] = "-1" # Force TF to use only the CPU
```

Share Edit Follow

answered Jul 26 '19 at 9:51



Edgar H

982 1 15 27



I have got the same error,lucky,I have got it fixed. my error is: the last time,I open the tensorflow `sess = tf.Session()` ,but I forgot close the session.

2

so I open the terminal, type command:



```
ps -aux | grep program_name
```

find the PID,and type command kill the PID:

```
kill -9 PID
```

Ok,the GPU is realase.

Share Edit Follow

edited Sep 4 '18 at 9:03



Mickael Maison

17.1k 7 44 46

answered Sep 4 '18 at 8:40



xiaxia wang

21 1



For me, closing and restarting my processes using Python worked.

2

I tried a few things here but they didn't work. For example, the



```
os.environ["CUDA_VISIBLE_DEVICES"] = "-1"
```

line. I think it is because I am using a newer version of Keras and Tensorflow. A lot of what I have read on the internet, include the official Keras tutorial, does not work because of version conflicts.

But I saw a couple of posts about more than one Python process running. So I shut down Jupyter, Anaconda, and PyCharm, and restarted everything. And then the error went away. It may or may not be what fixes is for you, but it is worth a try.

Share Edit Follow

answered Sep 11 '20 at 11:13



[Paul J Abernathy](#)

911 2 11 24



2

For me it was a runaway ipynb script which I thought I had terminated but was actually still running, thus my GPU was in use and this error appeared



Share Edit Follow

answered May 13 '20 at 0:53



[jarrettyeo](#)

1,539 1 17 31



1

I was stuck with this for days and I was able to get rid of this error eventually.

I had the wrong versions of tensorflow and cuda installed in my PC. Just make sure you have the right versions of tf,cuda and cudnn installed.



<https://i.stack.imgur.com/Laiii.png>

Use this link for reference.

Share Edit Follow

answered Sep 24 '19 at 17:14



[envi z](#)

197 2 11



0

I was using Jupyter Lab, but must have had the GPU locked from another TensorFlow program that I ran earlier. After killing Jupyter Lab and restarting it, the error went away.



Share Edit Follow

answered Oct 15 '19 at 11:52



[Nikhil Gupta](#)

641 7 12



0

Had python open in Dos window open in Windows 10. When running from my IDE, it gave the above message. Exiting from that Dos instance of python allowed me to get past this error when running from my IDE.





Share Edit Follow

answered 2 days ago



demongolem

8,551 13 81 99



I had the same error in Win10 using PyCharm on Keras 2.4.3 and TensorFlow 2.3.0

0



It seems like an error related to TensorFlow itself running on Windows. Solved the issue by closing PyCharm and rerun it as administrator.



Share Edit Follow

answered Dec 20 '20 at 12:27



pleasant

85 2 7



I had a similar kind of error while inferencing on a [Tensorflow model](#), Fixed that issue by **downgrading Tensorflow from 2.1 to 1.14**. Initially, I checked the GPU usage, it took all the GPU memory and couldn't able to perform the inference and found the following exception:

0



InternalError: 2 root error(s) found.

(0) Internal: Blas GEMM launch failed : a.shape=(86494, 257), b.shape=(257, 64), m=86494, n=64, k=257 [{{node log\_mel\_features/MatMul}}]

(1) Internal: Blas GEMM launch failed : a.shape=(86494, 257), b.shape=(257, 64), m=86494, n=64, k=257 [{{node log\_mel\_features/MatMul}}]

[[log\_mel\_features/Log/\_769]] Bellow are the commands I used:

```
pip uninstall tensorflow-gpu
pip install tensorflow-gpu==1.14
```

Share Edit Follow

answered May 5 '20 at 7:13



Akib Khan

31 4



Adding the following lines after imports solved the problem:

0



```
configuration = tf.compat.v1.ConfigProto()
configuration.gpu_options.allow_growth = True
session = tf.compat.v1.Session(config=configuration)
```



Share Edit Follow

edited Dec 31 '20 at 17:23



seuniojqap

2,501 8 28 36

answered Dec 31 '20 at 16:56



burhan rashid

31 5





0

Had the same error (Win10 using Keras and Visual Studio Code). Seems like TensorFlow was still active somehow even after terminating my script. Simply closing VS Code and restarting solved the issue.

[Share](#) [Edit](#) [Follow](#)

answered May 6 '19 at 16:57

[Peter](#)**1,177** 1 8 22

0

Try running the sample program simpleCUBLAS (it comes with CUDA) to test your CUBLAS installation and see if it works.



In my case (I am using Ubuntu) I had to reinstall CUDA to solve this issue. After I did that, simpleCUBLAS passed the test.



For some reason I started running into the same issue after a while, and I found that cleaning the directory .nv (inside my home folder) resolved the issue, and simpleCUBLAS test passed again.

[Share](#) [Edit](#) [Follow](#)

answered Jun 16 '17 at 23:30

[Tairone](#)**91** 1 6

I'm running Ubuntu 16.04 and simpleCUBLAS passes, but still getting the error "Blas GEMM launch failed" – [naisanza](#) Sep 5 '17 at 19:52

- 1 [@naisanza](#) you could be running two programs at once, both wanting to use tensorflow gpu – [kRazzy R](#) Feb 21 '18 at 3:37



**Highly active question.** Earn 10 reputation in order to answer this question. The reputation requirement helps protect this question from spam and non-answer activity.